

# Homework No. 5: Matrix Multiplication on a GPU Using CUDA

By Shahzad Khan

The goal for this assignment was to implement our indexing version of my matrix multiplication code to be run on a gpu. As a reminder our original indexing code was the following

```
bool multiMatrix(float* mat3, float* mat2, float* mat1, long nrows, long ncols){

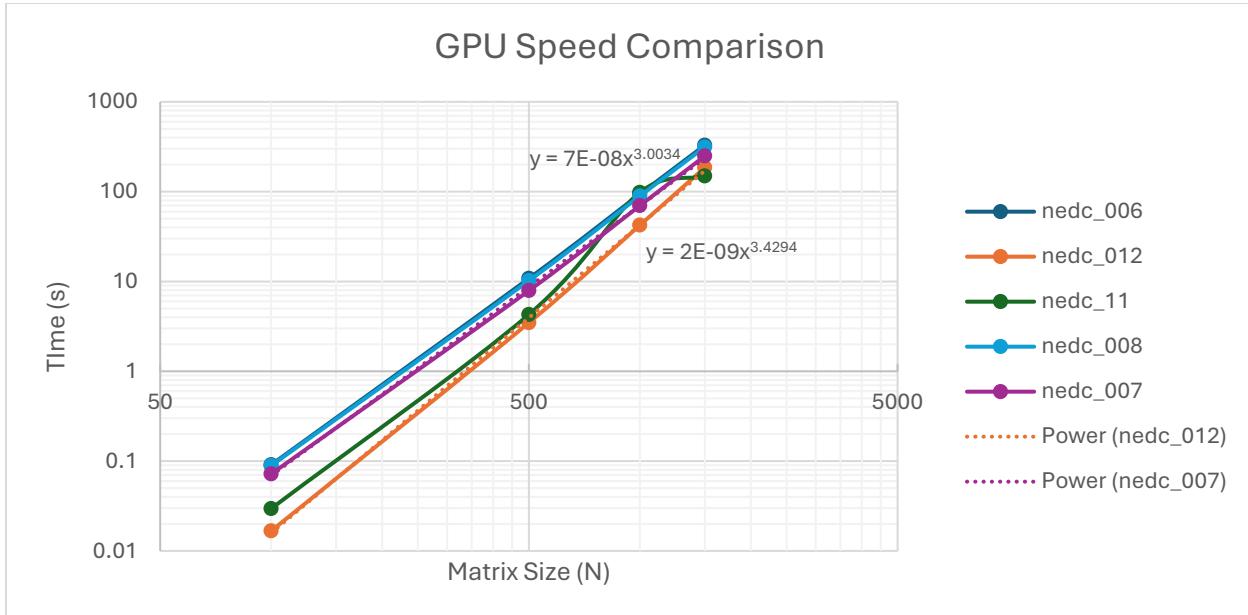
    if(!mat1 || !mat2 || !mat3){
        return false;
    }

    for (long i = 0; i < nrows; i++) {
        for (long j = 0; j < nrows; j++) {
            float sum = 0;
            for (long k = 0; k < ncols; k++) {
                sum += mat1[i * ncols + k] * mat2[k * nrows + j];
            }
            mat3[i * nrows + j] = sum;
        }
    }
    return true;
}
```

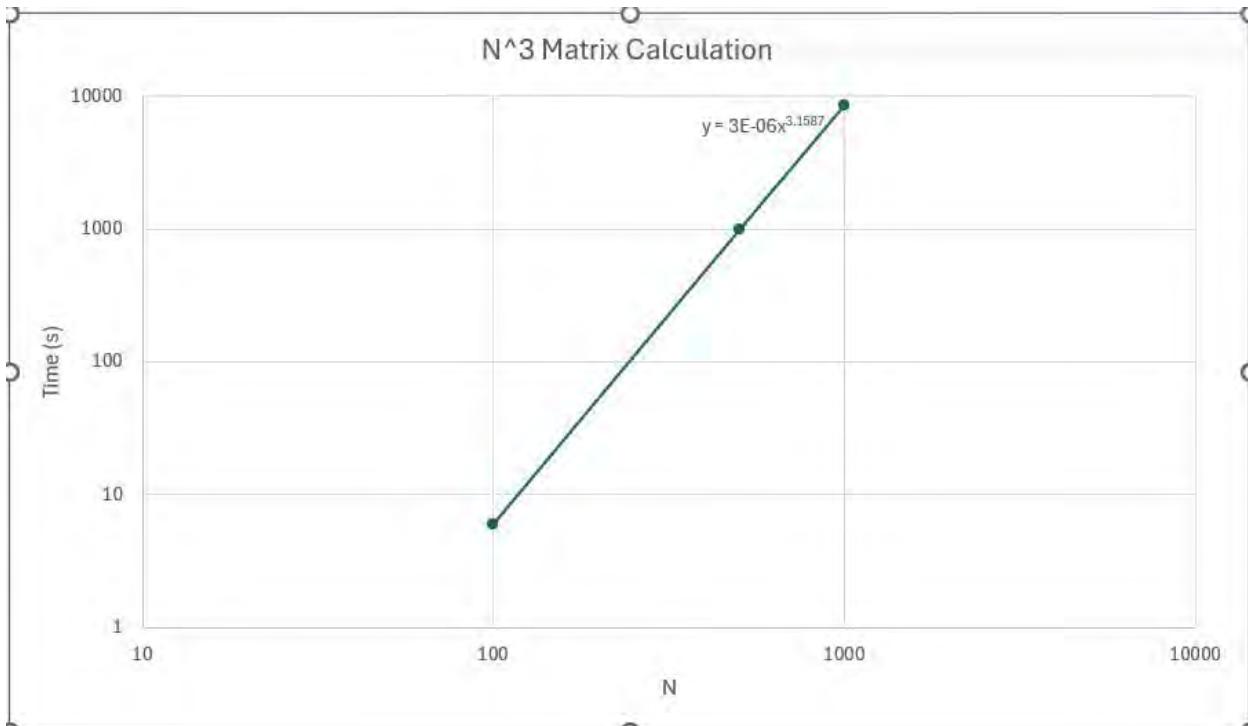
This gave us the time complexity of  $O(n^3)$  which when running this program with larger matrices caused our runtime to grow exponentially.

Now converting this code to run on a gpu was simple, since our goal was to simply over implement this only using 1 thread and 1 block `<<<1, 1>>>` to get a baseline for our next homework.

Our results are shown here



As we can see, it seems the fastest gpu was nedc\_012, but something interesting is that nedc\_11 did the best with my highest N, 1500. Even outclassing nedc\_012.



Looking back at my original hw\_01 graph we can see that the time complexity around the same at  $O(N^3)$  which makes sense because we aren't utilizing any parallelization

In conclusion, this assignment provided a good baseline for future GPU coding assignments with the different speeds of the GPUs. Nedc\_012 being overall the fastest and nedc\_011 having dip with a 1000x1000 matrix, and then being the fastest with a 1500x1500 matrix.